

## Section 4: Dijkstra's Algorithm

Amy Zhao

February 16, 2024

### 1 Dijkstra's Algorithm

At a high level, Dijkstra's algorithm is an algorithm which finds the shortest path between a starting vertex and another vertex in the graph. During each iteration, we expand the search by identifying the closest not-yet-processed vertex from the starting vertex.

```
Input: Graph  $G = (V, E)$  given as an adjacency list, starting vertex  $s$   
Post : Shortest distances from  $s$  to  $v \in V$ , stored in  $\text{dist}[v]$   
  
1 Function Dijkstra( $G = (V, E), s$ ):  
2    $Q \leftarrow$  new priority queue  
3   for  $v \in V$  do  
4      $\text{distance}[v] \leftarrow \infty$   
5      $\text{previous}[v] \leftarrow \text{null}$   
6     if  $v \neq s$  then  
7       add  $v$  with its distance to  $Q$   
8   end  
9    $\text{distance}[s] \leftarrow 0$   
10  while  $Q$  is not empty do  
11     $u \leftarrow$  vertex in  $Q$  with minimum distance  
12    for each unvisited neighbor  $v$  of  $u$  do  
13       $\text{temp} \leftarrow \text{distance}[u] + \text{weight}[u, v]$   
14      if  $\text{temp} < \text{distance}[v]$  then  
15         $\text{distance}[v] \leftarrow \text{temp}$   
16         $\text{previous}[v] \leftarrow u$   
17    end  
18  end
```


## 2 What Can Go Wrong?

### 2.1 Negative Edges

In this section, we illustrate a few examples when performing Dijkstra's will yield the wrong answer. First, suppose we have a graph  $G = (V, E)$  with edge weights  $L = \{\ell_e : e \in E\}$ , and we allow for the possibility of negative edges. Assume there are no negative cycles. Evidently, we know that Dijkstra's algorithm doesn't work on graphs with negative edges. Consider a modification to a graph with at least one negative edge:

- Find the minimum edge weight in  $L$ , denoted by  $\ell^* < 0$ .
- For all edges, change the edge weight to  $\ell'_e = \ell_e + |\ell^*| + 1$ , so that all of the new edge weights are positive. Denote the new set of edges as  $L' = \{\ell'_e : e \in E\}$ .
- Run Dijkstra's algorithm as usual with the new set of edge weights.

Will this change yield the shortest paths from a source vertex to all other vertices? Can we find a simple example or counterexample first?



**Figure 2.1:** Example of a single graph where Dijkstra's algorithm will not yield the correct answer on the modified graph. The shortest path in the original graph from  $a$  to  $c$  is  $a \rightarrow b \rightarrow c$ , while the shortest path in the modified graph is  $a \rightarrow c$ .

We claim that the proposed modification does not find the shortest path and provide a proof. For notation purposes, let  $u$  be the source vertex, and let  $v \in V$  be any other vertex in the graph. Let  $W(P)$  be the sum of edge weights on a path  $P$  using the edge weights from the original graph, and let  $W'(P)$  be the sum of the edge weights on a path  $P$  using the modified weights:

$$W(P) = \sum_{e \in P} \ell_e \qquad W'(P) = \sum_{e \in P} \ell'_e.$$

We set  $P$  to be the shortest path using the original edge weights. First, consider a path  $P$  containing  $k$  edges and the resulting sum on the modified edges:

$$\begin{aligned} W'(P) &= \sum_{e \in P} \ell'_e = \sum_{e \in P} (\ell_e + |\ell^*| + 1) \\ &= \left( \sum_{e \in P} \ell_e \right) + k(|\ell^*| + 1) \\ &= W(P) + k(|\ell^*| + 1). \end{aligned}$$

Consider a second path  $Q$  from  $u$  to  $v$  with  $j < k$  edges but greater weight  $W(Q) > W(P)$ . We write the sum  $W'(Q)$  as:

$$W'(Q) = \sum_{e \in Q} \ell'_e = W(Q) + j(|\ell^*| + 1).$$

These relationships yield the following implication:

$$\begin{aligned} W(Q) - W(P) &< (k - j)(|\ell^*| + 1) \\ \implies W'(Q) &= W(Q) + j(|\ell^*| + 1) < W(P) + k(|\ell^*| + 1) = W'(P). \end{aligned}$$

In general,  $W(P) < W(Q)$  does not imply  $W'(P) \leq W'(Q)$ , and thus the modification does not always return the shortest path with respect to the original edge weights.

## 2.2 Adding a Constant

Now, consider a positive weighted graph, and assume we use Dijkstra's algorithm to compute a shortest path  $P$  between two vertices  $u$  and  $v$ . Will Dijkstra's algorithm yield the same path if we add a positive constant to every edge?

Suppose the original graph has two possible paths from  $s$  to  $t$ : one which includes 50 edges with  $\ell_e = 1$  and one which includes a single edge with  $\ell_e = 51$ . When we run Dijkstra's algorithm on this graph, the shortest path will be the first path because it has a total weight of 50, as compared to the second which has total weight of 51.

If we add a constant to each edge of the graph, say  $c = 50$ , the total weight of the first path will be  $50(1 + 50) = 2550$ , while the weight of the second path will be

$51 + 50 = 101$ . Thus, this clearly illustrates why Dijkstra's will not work on this modified graph to return the shortest path length of the original graph. Paths with a greater number of edges will have a greater amount of weight added:

$$\begin{aligned} W'(P) &= \sum_{e \in P} (c + \ell_e) = |P|(c) + \sum_{e \in P} \ell_e = |P|(c) + W(P), \\ W'(Q) &= \sum_{e \in Q} (c + \ell_e) = |Q|(c) + \sum_{e \in Q} \ell_e = |Q|(c) + W(Q). \end{aligned}$$

$|P|(c) \gg |Q|(c)$  if  $|P| \gg |Q|$ , and this is solely dependent on the number of edges in the path.

### 2.3 Scaling by a Constant

Taking the same setup as in the previous section, we decide to scale each edge by a constant positive factor, instead of adding a constant. Will Dijkstra's algorithm work on this modified graph to produce the shortest path from the original?

Let  $P$  be the shortest path from  $s$  to  $t$  in the original graph, and let  $Q$  be an alternate path from  $s$  to  $t$  in the original graph. Assume also that the number of edges in the two paths may be different. We denote the weight of the original path with  $W$  and the weight of the modified path with  $W'$ .

As  $P$  is the shortest path, we have  $W(P) < W(Q)$ . We'll look at the scaling modification:

$$\begin{aligned} W'(P) &= \sum_{e \in P} (c \cdot \ell_e) = c \cdot \sum_{e \in P} \ell_e = c \cdot W(P). \\ W'(Q) &= \sum_{e \in Q} (c \cdot \ell_e) = c \cdot \sum_{e \in Q} \ell_e = c \cdot W(Q). \end{aligned}$$

Thus, we have:

$$W'(P) = c \cdot W(P) < c \cdot W(Q) = W'(Q) \implies W'(P) < W'(Q),$$

which implies that we will obtain the same shortest path in the modified graph when we run Dijkstra's.

### 3 Example: Fuel Capacity

Suppose we have a set of cities connected by highways, given in the form of an undirected graph  $G = (V, E)$ . Each highway  $e \in E$  has a specified length  $\ell_e > 0$ . You are planning a road trip and want to determine if you can travel from city  $s$  to city  $t$  with a car that has a fuel capacity of  $L$  miles. You can refuel at each city, but not between, which means you can take a route if every one of its highways is smaller in length than  $L$ .

#### 3.1 Reachability

If we know the car's fuel capacity, how do we determine if we can reach  $t$  from  $s$ ? Revisiting the previous discussion section's examples, we can consider modifying the graph by taking out any edges that are longer than  $L$ . Running DFS on the resulting graph will check reachability between  $s$  and  $t$ , yielding the correct answer. Modification and DFS are  $O(|V| + |E|)$  runtime.

#### 3.2 Determining Needed Capacity

However, we face a trickier problem if we don't know the car's capacity. Suppose we are purchasing a new vehicle and want to ensure we buy a car with enough fuel capacity to travel from  $s$  to  $t$ . Notably, we can refuel at each city as needed, so the car's capacity need only to be at least as large as the longest edge of the smallest path from  $s$  to  $t$ . In other words, the fuel capacity needed is constrained by the  $s$  to  $t$  path whose longest edge is the smallest; we'll call this the smallest max length path.

To determine this length, we adapt Dijkstra's algorithm. At each vertex  $u$ , we track the constraining edge, instead of the distance, in  $m[u]$ . Every time we dequeue a vertex  $u$ , we compare the length of the edge  $e = \{u, v\}$  with the current smallest max length:

- $\max\{\ell_e, m[u]\} < m[v]$ : set  $m[v] = \max\{\ell_e, m[u]\}$  and rebalance the priority queue; do nothing otherwise.

The algorithm returns the value at  $m[t]$ , which will accurately report the constraining edge in the smallest path.

	<b>Input</b> : Graph $G = (V, E)$ given as an adjacency list with weights $\ell_e > 0$ for $e \in E$ , starting vertex $s$ , target vertex $t$
	<b>Output</b> : Smallest max length path
1	<b>Function</b> FuelConstraint( $G = (V, E), s, t$ ):
2	<b>for</b> $v \in V$ <b>do</b>
3	<b>if</b> $v = s$ <b>then</b>
4	$m[v] \leftarrow 0$
5	<b>else</b>
6	$m[v] \leftarrow \infty$
7	<b>end</b>
8	$Q \leftarrow$ new priority queue with objects $v \in V$ , associated with priority $m[v]$
9	<b>while</b> $Q$ is not empty <b>do</b>
10	$v \leftarrow \text{ExtractMin}(Q)$
11	<b>for</b> $w \in \text{Adj}[v]$ <b>do</b>
12	<b>if</b> $m[w] > \max\{\ell_{vw}, m[v]\}$ <b>then</b>
13	$m[w] \leftarrow \max\{\ell_{vw}, m[v]\}$
14	Adjust priority key for $w$
15	<b>end</b>
16	<b>end</b>
17	<b>return</b> $m[t]$

Tracking the smallest max length path does not alter the runtime of Dijkstra's algorithm, so the overall time complexity remains  $O((|E| + |V|) \log(|V|))$ .

## 4 Example: Number of Shortest Paths

We know that Dijkstra's algorithm is used to find the shortest paths in a weighted undirected graph. Is it possible to augment the algorithm to count the number of shortest paths from  $s$  to  $t$ ? Recall that in running Dijkstra's algorithm, all other vertices with a shorter path from  $s$  will have already been explored before we expand vertex  $a$  in Dijkstra's algorithm. This means that we know the shortest path from  $s$  to these vertices is known at this point, and we can simply record additional information about the shortest paths based off a couple comparisons.

At some iteration of Dijkstra's algorithm, let's consider the expansion of vertex  $u$  on the edge  $\{u, v\}$  with length  $\ell_e$ :

- $\text{dist}[v] > \text{dist}[u] + \ell_e$ : We find a shorter path to  $v$ , so we set  $\text{dist}[v] = \text{dist}[u] + \ell_e$ , and we update the number of paths to  $v$  to be equal to the number of paths

from  $s$  to  $u$ ,  $\text{numpaths}[v] = \text{numpaths}[u]$ .

- $\text{dist}[v] = \text{dist}[u] + \ell_e$ : We find a path with the same distance as previously recorded, so there are no updates to make to the distance array. However, we increase the number of paths to  $v$ ,  $\text{numpaths}[v] += \text{numpaths}[u]$ , as we care about the path distances, rather than the actual edges of the paths.

The algorithm returns the value at  $\text{numpaths}[t]$ , which will accurately report the number of shortest paths because the value is updated in parallel with any updates to the distance array. This means that whenever  $v$  has the smallest distance value, there are no other unexplored paths from  $s \rightarrow v$  that could decrease the value we've stored at  $\text{dist}[v]$ .

<b>Input</b> : Graph $G = (V, E)$ given as an adjacency list, starting vertex $s$ , target vertex $t$ <b>Output</b> : Number of shortest paths from $s$ to $t$	<pre> 1 <b>Function</b> NumShortestPaths(<math>G = (V, E), s, t</math>): 2   <b>for</b> <math>v \in V</math> <b>do</b> 3     <math>\text{dist}[v] \leftarrow \infty</math> 4     <math>\text{numpaths}[v] \leftarrow 0</math> 5   <b>end</b> 6   <math>\text{dist}[s] \leftarrow 0</math> 7   <math>\text{numpaths}[s] \leftarrow 1</math> 8   <math>Q \leftarrow</math> new priority queue 9   <b>while</b> <math>Q</math> is not empty <b>do</b> 10    <math>v \leftarrow</math> vertex in <math>Q</math> with minimum distance 11    <b>for</b> <math>w \in \text{Adj}[v]</math> <b>do</b> 12      <b>if</b> <math>\text{dist}[w] &gt; \text{dist}[v] + \ell_{vw}</math> <b>then</b> 13        <math>\text{dist}[w] \leftarrow \text{dist}[v] + \ell_{vw}</math> 14        <math>\text{ChangeKey}(P, w)</math> 15        <math>\text{numpaths}[w] \leftarrow \text{numpaths}[v]</math> 16      <b>else</b> 17        <math>\text{numpaths}[w] \leftarrow \text{numpaths}[w] + \text{numpaths}[v]</math> 18      <b>end</b> 19    <b>end</b> 20  <b>return</b> <math>\text{numpaths}[t]</math> </pre>
---	---

The runtime of the algorithm is the same as Dijkstra's,  $O((|E| + |V|) \log |V|)$ , as the comparisons and updates to the  $\text{numpaths}$  array don't impose any additional dominating runtime.